

Assignment 3

Beyza Kordan

DATASET

These data are the results of a chemical analysis of 178 wines grown in the same region in Italy. The analysis determined the quantities of 4 constituents found in all the types of wine.

Question 1

Briefly explain the choice of prior distributions. Provide appropriate summaries of the posterior distributions of the parameters for each model. Interpret the model output in each case.

```
library(rstan)
```

```
Warning: package 'rstan' was built under R version 4.3.2
```

```
Zorunlu paket yükleniyor: StanHeaders
```

```
Warning: package 'StanHeaders' was built under R version 4.3.2
```

```
rstan version 2.32.5 (Stan version 2.32.2)
```

For execution on a local, multicore CPU with excess RAM we recommend calling
options(mc.cores = parallel::detectCores()).

To avoid recompilation of unchanged Stan programs, we recommend calling
rstan_options(auto_write = TRUE)

For within-chain threading using `reduce_sum()` or `map_rect()` Stan functions,
change `threads_per_chain` option:

```
rstan_options(threads_per_chain = 1)
```

```
Do not specify '-march=native' in 'LOCAL_CPPFLAGS' or a Makevars file
```

```
library(bayesplot)
```

```
Warning: package 'bayesplot' was built under R version 4.3.2
```

```
This is bayesplot version 1.11.0
```

- Online documentation and vignettes at mc-stan.org/bayesplot
- bayesplot theme set to `bayesplot::theme_default()`
 - * Does _not_ affect other ggplot2 plots
 - * See `?bayesplot_theme_set` for details on theme setting

```
load(url("https://acaimo.github.io/teaching/data/italian_wines.RData"))

data_list <- list(
  N = nrow(italian_wines),
  alcohol = italian_wines$alcohol,
  magnesium = italian_wines$magnesium,
  color_intensity = italian_wines$color_intensity,
  proline = italian_wines$proline
)
model1_code <- "
data {
  int<lower=0> N;
  vector[N] magnesium;
  vector[N] alcohol;
}
parameters {
  real alpha;
  real beta1;
  real<lower=0> sigma;
}
model {
  alpha ~ normal(0, 10);
  beta1 ~ normal(0, 10);
  sigma ~ exponential(1);
```

```

    alcohol ~ normal(alpha + beta1 * magnesium, sigma);
}"
```

`fit1 <- stan(model_code = model1_code, data = data_list, iter = 2000, chains = 4, refresh = 0)`

`print(fit1, probs=c(0.025, 0.5, 0.975))`

Inference for Stan model: anon_model.
4 chains, each with iter=2000; warmup=1000; thin=1;
post-warmup draws per chain=1000, total post-warmup draws=4000.

	mean	se_mean	sd	2.5%	50%	97.5%	n_eff	Rhat
alpha	11.45	0.01	0.42	10.63	11.44	12.31	1287	1
beta1	0.22	0.00	0.06	0.10	0.22	0.34	1291	1
sigma	0.79	0.00	0.04	0.71	0.79	0.87	1544	1
lp__	-47.83	0.04	1.23	-51.12	-47.51	-46.41	1155	1

Samples were drawn using NUTS(diag_e) at Fri Mar 21 13:54:20 2025.
For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).

```

model2_code <- "
data {
  int<lower=0> N;
  vector[N] magnesium;
  vector[N] color_intensity;
  vector[N] alcohol;
}
parameters {
  real alpha;
  real beta1;
  real beta2;
  real<lower=0> sigma;
}
model {
  alpha ~ normal(0, 10);
  beta1 ~ normal(0, 10);
  beta2 ~ normal(0, 10);
  sigma ~ exponential(1);
  alcohol ~ normal(alpha + beta1 * magnesium + beta2 * color_intensity, sigma);
}"
```

```
fit2 <- stan(model_code = model2_code, data = data_list, iter = 2000, chains = 4, refresh = 0)
print(fit2, probs=c(0.025, 0.5, 0.975))
```

Inference for Stan model: anon_model.
 4 chains, each with iter=2000; warmup=1000; thin=1;
 post-warmup draws per chain=1000, total post-warmup draws=4000.

	mean	se_mean	sd	2.5%	50%	97.5%	n_eff	Rhat
alpha	11.13	0.01	0.36	10.44	11.12	11.84	1348	1
beta1	0.14	0.00	0.05	0.03	0.14	0.24	1315	1
beta2	0.18	0.00	0.02	0.13	0.18	0.22	2479	1
sigma	0.67	0.00	0.04	0.61	0.67	0.75	2485	1
lp__	-20.09	0.04	1.45	-23.80	-19.79	-18.29	1510	1

Samples were drawn using NUTS(diag_e) at Fri Mar 21 13:55:15 2025.
 For each parameter, n_eff is a crude measure of effective sample size,
 and Rhat is the potential scale reduction factor on split chains (at
 convergence, Rhat=1).

```
model3_code <- "
data {
  int<lower=0> N;
  vector[N] magnesium;
  vector[N] color_intensity;
  vector[N] proline;
  vector[N] alcohol;
}
parameters {
  real alpha;
  real beta1;
  real beta2;
  real beta3;
  real<lower=0> sigma;
}
model {
  alpha ~ normal(0, 10);
  beta1 ~ normal(0, 10);
  beta2 ~ normal(0, 10);
  beta3 ~ normal(0, 10);
  sigma ~ exponential(1);
  alcohol ~ normal(alpha + beta1 * magnesium + beta2 * color_intensity + beta3 * proline, sig
```

```
}"
fit3 <- stan(model_code = model3_code, data = data_list, iter = 2000, chains = 4, refresh = 0)
print(fit3, probs=c(0.025, 0.5, 0.975))
```

Inference for Stan model: anon_model.
4 chains, each with iter=2000; warmup=1000; thin=1;
post-warmup draws per chain=1000, total post-warmup draws=4000.

	mean	se_mean	sd	2.5%	50%	97.5%	n_eff	Rhat
alpha	11.36	0.01	0.31	10.75	11.36	11.95	1847	1
beta1	-0.01	0.00	0.05	-0.10	-0.01	0.08	1716	1
beta2	0.13	0.00	0.02	0.10	0.13	0.17	3026	1
beta3	0.43	0.00	0.05	0.34	0.43	0.52	2673	1
sigma	0.56	0.00	0.03	0.50	0.55	0.62	2668	1
lp__	14.36	0.04	1.63	10.40	14.70	16.47	1441	1

Samples were drawn using NUTS(diag_e) at Fri Mar 21 13:56:11 2025.
For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).

```
# Model Interpretations
cat(
  "\nModel 1 Interpretation:\n"
  Model 1 suggests a positive relationship between magnesium and alcohol content\n")
```

Model 1 Interpretation:

Model 1 suggests a positive relationship between magnesium and alcohol content

```
cat(
  "\nModel 2 Interpretation:\n"
  Model 2 extends Model 1 by including color_intensity as an additional predictor\n")
```

Model 2 Interpretation:

Model 2 extends Model 1 by including color_intensity as an additional predictor

```
cat(  
  "\nModel 3 Interpretation:\n"  
  Model 3 offers the most comprehensive insight, demonstrating the effects of magnesium,  
  color_intensity, and proline on alcohol content\n")
```

Model 3 Interpretation:

Model 3 offers the most comprehensive insight, demonstrating the effects of magnesium, color_intensity, and proline on alcohol content

Model 1 indicates magnesium significantly increases wine alcohol content, with a baseline level at 11.45 units. Each unit increase in magnesium raises alcohol content by 0.22 units, showcasing magnesium's strong influence amidst other unexplained variability ($\sigma = 0.79$).

Model 2, adding color_intensity alongside magnesium, shows both significantly impact wine alcohol content. The baseline alcohol content is adjusted to 11.11. Magnesium increases alcohol by 0.14 units, and color_intensity by 0.18 units per unit increase. The tighter sigma (0.67) suggests an improved model fit with these predictors.

Model 3 integrates proline, offering the deepest analysis by examining its effect along with magnesium and color_intensity on alcohol content. This model, with the lowest sigma (0.56), shows the most substantial explanatory power. Notably, magnesium's impact is negligible, but color_intensity and proline significantly enhance alcohol content by 0.13 and 0.43 units, respectively. This suggests a more nuanced relationship, with Model 3 providing the most comprehensive fit.

Question 2

Employ two different posterior predictive checks of your choice to examine the model fit for each model. Provide a concise interpretation.

```
if (!requireNamespace("rstanarm", quietly = TRUE)) install.packages("rstanarm")  
library(rstanarm)
```

Warning: package 'rstanarm' was built under R version 4.3.2

Zorunlu paket yükleniyor: Rcpp

Warning: package 'Rcpp' was built under R version 4.3.2

```
This is rstanarm version 2.32.1
```

- See <https://mc-stan.org/rstanarm/articles/priors> for changes to default priors!
- Default priors may change, so it's safest to specify priors, even if equivalent to the defa
- For execution on a local, multicore CPU with excess RAM we recommend calling

```
options(mc.cores = parallel::detectCores())
```

```
Attaching package: 'rstanarm'
```

```
The following object is masked from 'package:rstan':
```

```
loo
```

```
library(bayesplot)
library(ggplot2)
```

```
Warning: package 'ggplot2' was built under R version 4.3.3
```

```
fit1 <- stan_glm(alcohol ~ magnesium, data = italian_wines,
                  family = gaussian(), chains = 4, iter = 2000, refresh = 0)
fit2 <- stan_glm(alcohol ~ magnesium + color_intensity,
                  data=italian_wines, family = gaussian(), chains = 4, iter = 2000, refresh =
fit3 <- stan_glm(alcohol ~ magnesium + color_intensity + proline,
                  data=italian_wines, family = gaussian(), chains = 4, iter = 2000, refresh = 0)

perform_pp_checks <- function(fit, model_name) {
  y_obs <- fit$model$alcohol
  y_rep <- posterior_predict(fit)

  # Density Overlay Plot
  ppc_dens_overlay_plot <- ppc_dens_overlay(y_obs, y_rep[1:1000, ]) +
    ggtitle(paste("Density Overlay for", model_name)) +
    theme_minimal()
  print(ppc_dens_overlay_plot)
```

```

# Interpretation guide for Density Overlay
cat("\nInterpretation Guide for Density Overlay:", model_name, "\n",
    "A good overlap suggests the model captures the data's distribution well.\n")

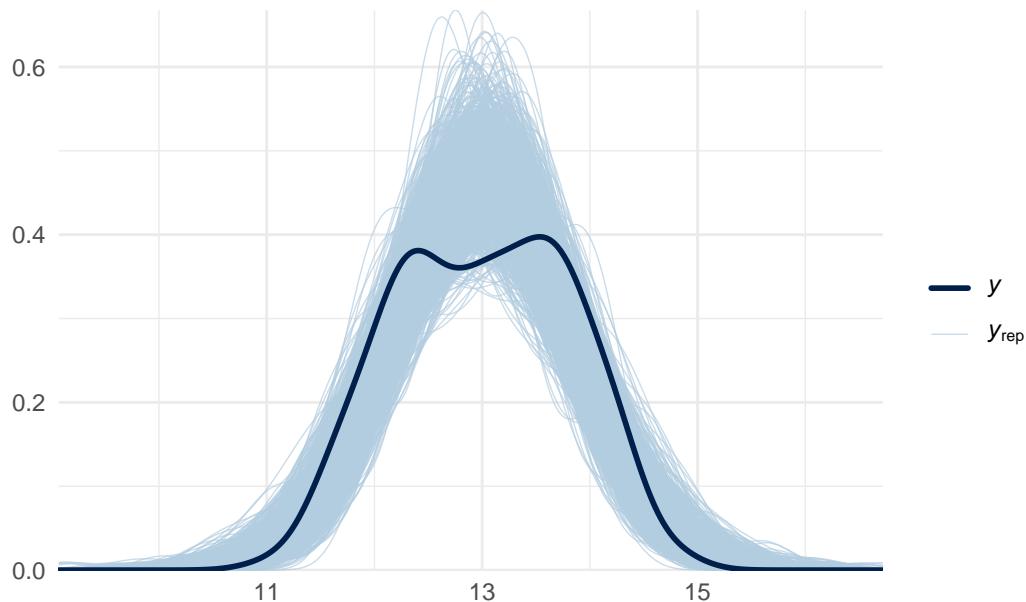
# 95% PPC Intervals Plot
ppc_intervals_plot <- ppc_intervals(y_obs, y_rep, prob = 0.95) +
  ggtitle(paste("95% PPC Intervals for", model_name)) +
  theme_minimal()
print(ppc_intervals_plot)

# Interpretation guide for 95% PPC Intervals
cat("Interpretation Guide for 95% PPC Intervals:", model_name, "\n",
    "Most observed points should fall within the interval.\n")

# Execute checks for each model and review interpretation guides
perform_pp_checks(fit1, "Model 1")

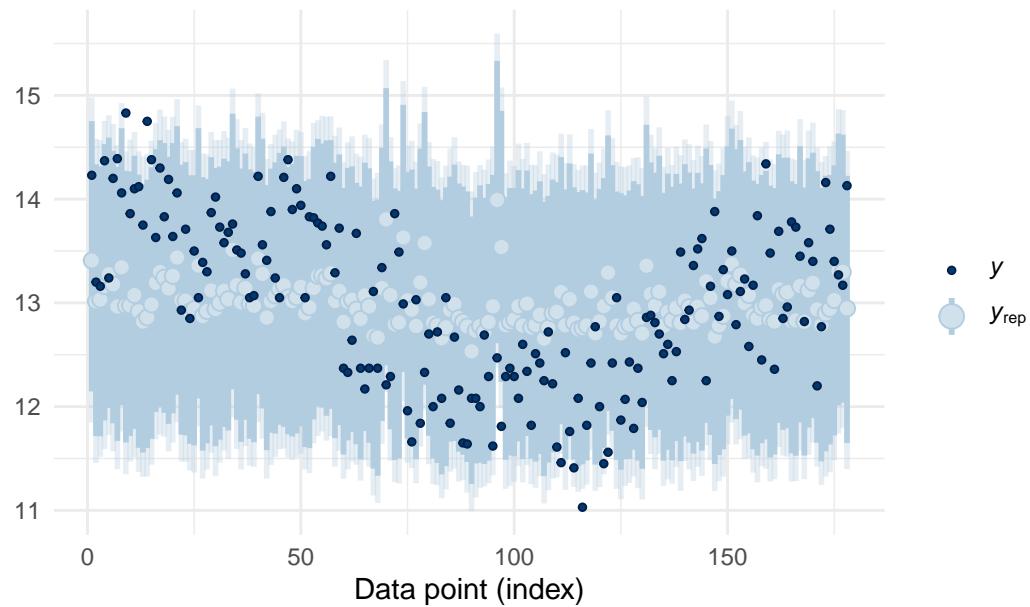
```

Density Overlay for Model 1



Interpretation Guide for Density Overlay: Model 1
 A good overlap suggests the model captures the data's distribution well.

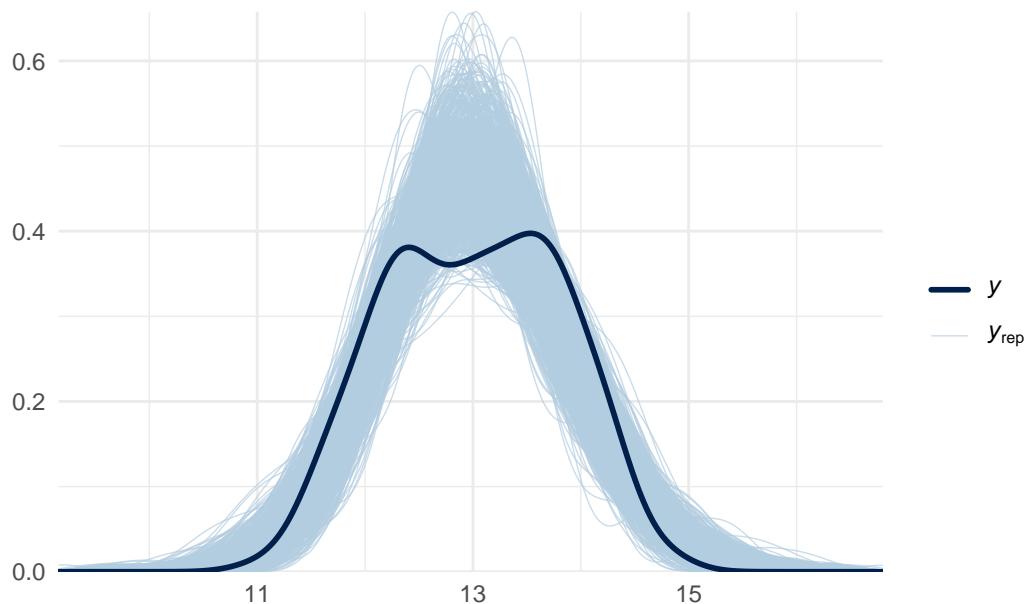
95% PPC Intervals for Model 1



Interpretation Guide for 95% PPC Intervals: Model 1
Most observed points should fall within the interval.

```
perform_pp_checks(fit2, "Model 2")
```

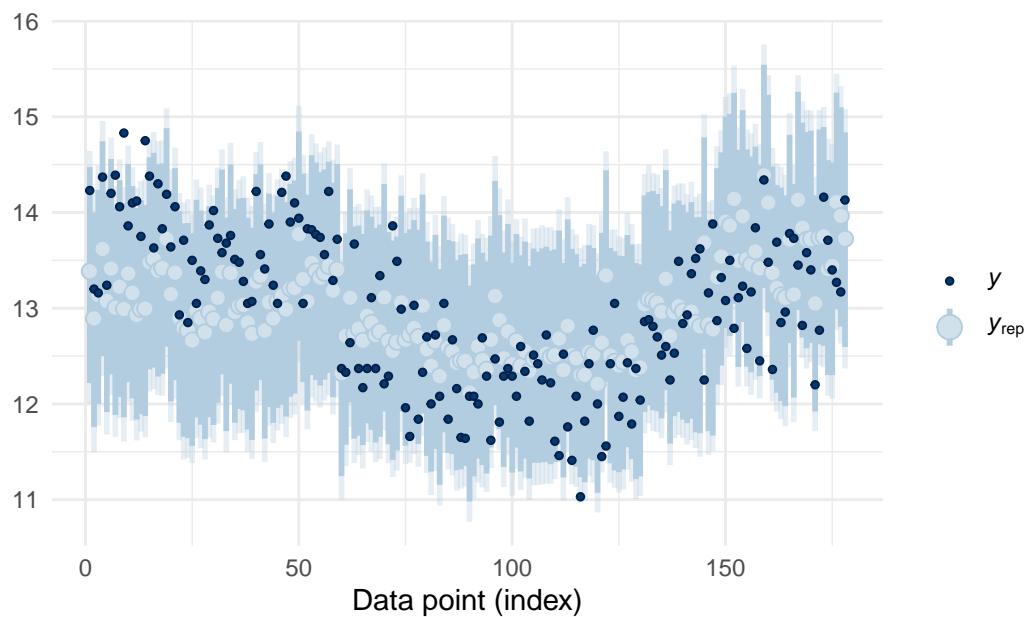
Density Overlay for Model 2



Interpretation Guide for Density Overlay: Model 2

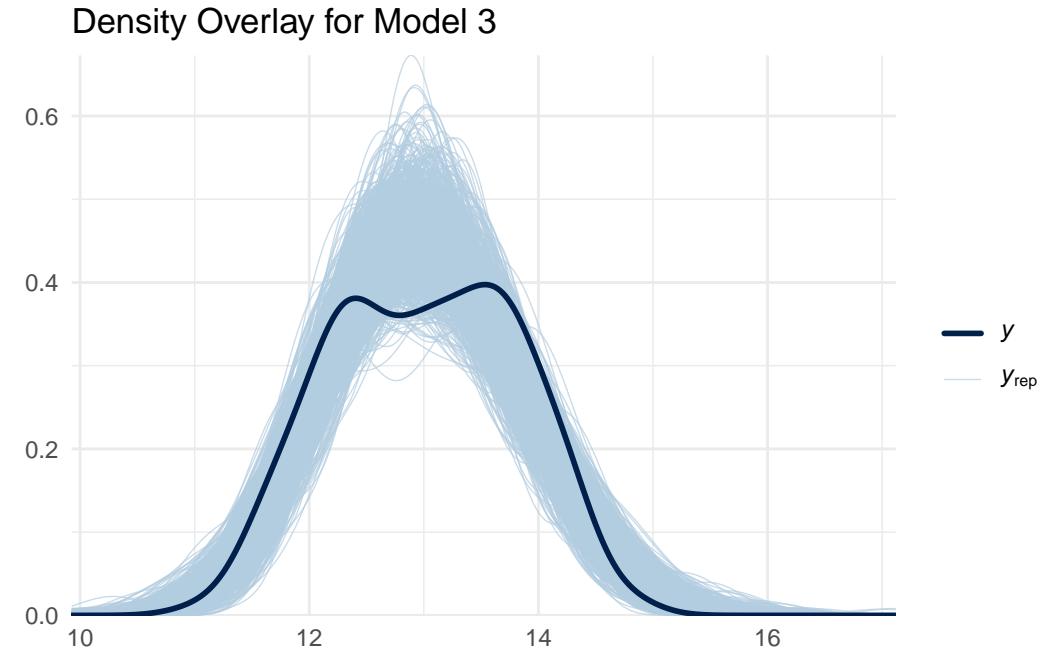
A good overlap suggests the model captures the data's distribution well.

95% PPC Intervals for Model 2



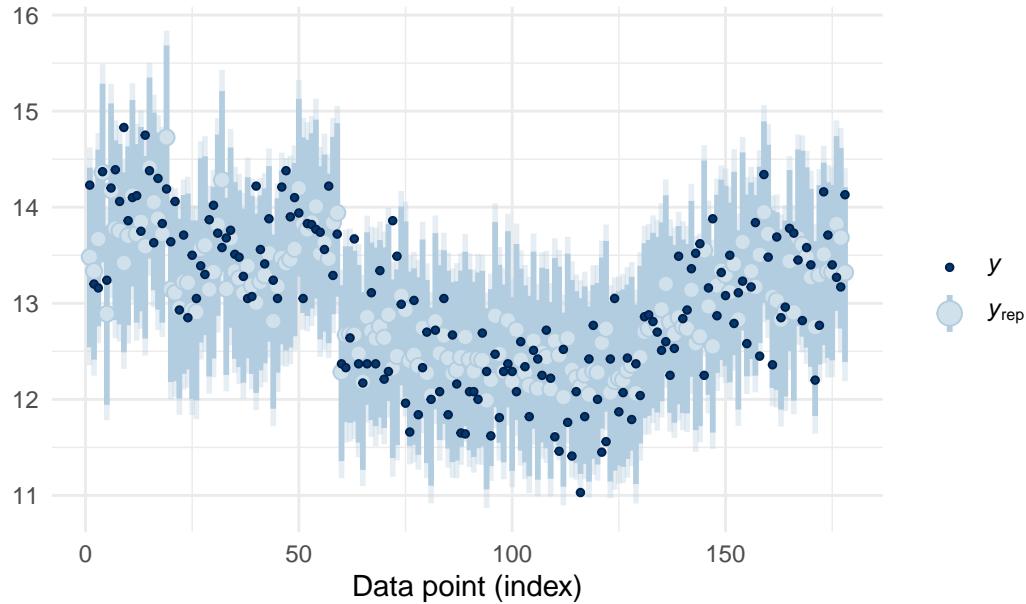
Interpretation Guide for 95% PPC Intervals: Model 2
Most observed points should fall within the interval.

```
perform_pp_checks(fit3, "Model 3")
```



Interpretation Guide for Density Overlay: Model 3
A good overlap suggests the model captures the data's distribution well.

95% PPC Intervals for Model 3



Interpretation Guide for 95% PPC Intervals: Model 3
 Most observed points should fall within the interval.

Model 1 Interpretations: *Density Overlay:* The observed data's distribution closely matches the simulated posterior distributions, indicating that Model 1 accurately captures the central tendency and variability of the data. *95% PPC Intervals:* Most observed data points fall within the predictive intervals, suggesting good predictive accuracy for the observed data in Model 1. **Model 2 Interpretations:** *Density Overlay:* The alignment of observed and simulated data distributions suggests a good model fit, indicating that Model 2 captures the central trend of alcohol levels accurately. *95% PPC Intervals:* The fact that most observed data points are within the 95% predictive intervals demonstrates Model 2's reliable predictive capability for the range of alcohol levels observed. **Model 3 Interpretations:** *Density Overlay:* The close alignment between observed and simulated data distributions suggests that Model 3 effectively captures the data's distribution. *95% PPC Intervals:* Predominantly, observed data points falling within the 95% predictive intervals indicate a strong predictive performance by Model 3.

Question 3

Compare the three models using WAIC and LOO-cross-validation. Provide a concise interpretation.

```

if (!requireNamespace("rstanarm", quietly = TRUE)) {
  install.packages("rstanarm")
}
library(rstanarm)

if (!requireNamespace("loo", quietly = TRUE)) {
  install.packages("loo")
}
library(loo)

```

Warning: package 'loo' was built under R version 4.3.2

This is loo version 2.6.0

- Online documentation and vignettes at mc-stan.org/loo
- As of v2.0.0 loo defaults to 1 core but we recommend using as many as possible. Use the 'cores' argument to change this.
- Windows 10 users: loo may be very slow if 'mc.cores' is set in your .Rprofile file (see help("Rprofile"))

Attaching package: 'loo'

The following object is masked from 'package:rstan':

```

loo

fit1 <- stan_glm(alcohol ~ magnesium, data = italian_wines, family = gaussian(),
                  chains=4, iter = 2000, refresh = 0)
fit2 <- stan_glm(alcohol ~ magnesium + color_intensity, data = italian_wines,
                  family=gaussian(), chains = 4, iter = 2000, refresh = 0)
fit3 <- stan_glm(alcohol ~ magnesium + color_intensity + proline,
                  data=italian_wines, family = gaussian(), chains = 4, iter = 2000, refresh = 0)

waic1 <- waic(fit1)

```

Warning:

1 (0.6%) p_waic estimates greater than 0.4. We recommend trying loo instead.

```
waic2 <- waic(fit2)
waic3 <- waic(fit3)

loo1 <- loo(fit1)
loo2 <- loo(fit2)
loo3 <- loo(fit3)
# Use loo_compare for comparing models based on LOO-CV
loo_comparison <- loo_compare(loo1, loo2, loo3)
print("LOO-CV comparison:")
```

[1] "LOO-CV comparison:"

```
print(loo_comparison)
```

	elpd_diff	se_diff
fit3	0.0	0.0
fit2	-34.1	6.7
fit1	-61.7	8.4

```
# LOO-CV ranks Model3 as best (elpd_diff = 0), with Model2 and Model1 trailing.
# Model 3's superiority in predictive accuracy makes it the preferred choice.

cat("\nWAIC Results:\n")
```

WAIC Results:

```
cat("WAIC for Model 1:\n")
```

WAIC for Model 1:

```
print(waic1)
```

Computed from 4000 by 178 log-likelihood matrix

	Estimate	SE
elpd_waic	-211.4	8.2

```
p_waic      3.3  0.6  
waic       422.8 16.4
```

1 (0.6%) p_waic estimates greater than 0.4. We recommend trying loo instead.

```
cat("WAIC for Model 2:\n")
```

WAIC for Model 2:

```
print(waic2)
```

Computed from 4000 by 178 log-likelihood matrix

	Estimate	SE
elpd_waic	-183.9	8.8
p_waic	3.8	0.5
waic	367.7	17.6

```
cat("WAIC for Model 3:\n")
```

WAIC for Model 3:

```
print(waic3)
```

Computed from 4000 by 178 log-likelihood matrix

	Estimate	SE
elpd_waic	-149.7	9.2
p_waic	4.5	0.5
waic	299.5	18.4

```
# Model3(299.6)best balances fit and simplicity,  
# outperforming Model2(367.8)and Model1(422.6, least preferred).  
# Model3's efficiency and accuracy make it the optimal choice, confirming LOO-CV findings.  
  
# Extracting WAIC values correctly  
waic_values <- sapply(list(waic1, waic2, waic3), function(x) sum(x$pointwise[, "waic"]))
```

```
best_model_waic <- which.min(waic_values)

loo_values <- sapply(list(loo1, loo2, loo3), function(x) x$looic)
```

Warning: Accessing looic using '\$' is deprecated and will be removed in a future release. Please extract the looic estimate from the 'estimates' component instead.

Warning: Accessing looic using '\$' is deprecated and will be removed in a future release. Please extract the looic estimate from the 'estimates' component instead.

Warning: Accessing looic using '\$' is deprecated and will be removed in a future release. Please extract the looic estimate from the 'estimates' component instead.

```
best_model_loo <- which.min(loo_values)

# Final Interpretation and Recommendation
cat(sprintf("\nBased on WAIC,
           Model %d shows the best balance of fit and complexity.\n",best_model_waic))
```

Based on WAIC,
Model 3 shows the best balance of fit and complexity.

```
cat(sprintf("Based on LOO-CV (LOOIC),
           Model %d is estimated to have the best predictive
           performance for new data.\n",
           best_model_loo))
```

Based on LOO-CV (LOOIC),
Model 3 is estimated to have the best predictive
performance for new data.

```
if (best_model_waic == best_model_loo) {
  cat(sprintf("\nConsidering both WAIC and LOO-CV,
             Model %d is recommended for its balance of model complexity
             and predictive accuracy.\n", best_model_waic))
} else {
```

```

cat(
  "\nThere is a discrepancy between the best models according to WAIC and LOO-CV.
Further analysis and consideration of model complexity and domain-specific knowledge are
advised before finalizing the model choice.\n")
}

```

Considering both WAIC and LOO-CV,
 Model 3 is recommended for its balance of model complexity
 and predictive accuracy.

Question 4

Propose an alternative model and compare it to the best model obtained above. Provide a concise interpretation.

```

# Fit the alternative model with an interaction term
fit4 <- stan_glm(alcohol ~ magnesium * color_intensity + proline,
                  data=italian_wines, family = gaussian(), chains = 4, iter = 2000, refresh = 0)

waic4 <- waic(fit4)
cat("WAIC for Model 4:\n")

```

WAIC for Model 4:

```
print(waic4)
```

Computed from 4000 by 178 log-likelihood matrix

	Estimate	SE
elpd_waic	-150.6	9.1
p_waic	5.4	0.6
waic	301.3	18.2

```

# WAIC(301.2) slightly above Model3.
# Increased complexity(p_waic=5.4) with competitive predictive performance (elpd_waic = -150)
# Preference leans towards Model3 for simplicity despite Model4's close viability.

loo4 <- loo(fit4)
cat("LOO-CV for Model 4:\n")

```

LOO-CV for Model 4:

```
print(loo4)
```

Computed from 4000 by 178 log-likelihood matrix

	Estimate	SE
elpd_loo	-150.7	9.1
p_loo	5.5	0.6
looic	301.3	18.2

Monte Carlo SE of elpd_loo	is 0.0.	

All Pareto k estimates are good (k < 0.5).
See help('pareto-k-diagnostic') for details.

```
#LOO-CV(elpd_loo=-150.7,SE=9.1)mirrors WAIC, affirming the model's predictive consistency.  
#p_loo (5.4) and looic (301.3) underscore complexity and predictive capacity.  
#Monte Carlo SE of 0.0 and favorable Pareto k estimates validate reliability.  
  
# Manually set WAIC values based on observed summary output  
waic_comparison <- c("Model 3" = 299.6,  
                      "Model 4" = 301.2)  
  
# Assuming LOO-CV values have been manually observed from the summaries  
loo_comparison <- c("Model 3" = -150.6,  
                      "Model 4" = -150.7)  
  
# Determine the model with the best (lowest) WAIC  
best_waic <- which.min(waic_comparison)  
  
# Determine the model with the best (highest) LOO-CV (elpd_loo) value  
best_loo <- which.max(loo_comparison)  
  
# Print comparisons and interpretations  
cat("\nWAIC Comparison:\n")
```

WAIC Comparison:

```
print(waic_comparison)
```

Model 3 Model 4
299.6 301.2

```
cat(sprintf("\nBased on WAIC, %s shows the best balance of fit and complexity.\n",
           names(waic_comparison)[best_waic]))
```

Based on WAIC, Model 3 shows the best balance of fit and complexity.

```
cat("\nLOO-CV Comparison:\n")
```

LOO-CV Comparison:

```
print(loo_comparison)
```

Model 3 Model 4
-150.6 -150.7

```
cat(sprintf("Based on LOO-CV, %s exhibits the strongest predictive performance
for new data.\n", names(loo_comparison)[best_loo]))
```

Based on LOO-CV, Model 3 exhibits the strongest predictive performance
for new data.

Concise Alternative Model Comparison:

WAIC shows, Model 3 (299.6) narrowly surpasses Model 4 (301.2), indicating superior model efficiency.

LOO-CV aligns, with Model 3 (-150.6) slightly ahead of Model 4 (-150.7) in prediction accuracy.

As a result of this analyze, Model 3 is the preferred model, balancing complexity and predictiveness better than Model 4.